

Eric Andrews, Christopher Chapman

CPS377: Database Systems, Professor Justin Brody

GitHub: <https://github.com/CWMChapman/DBMS-Implementation>

Presentation: 2-5pm December 1, 2020

EMPIRICALLY BENCHMARKING TREE-BASED AND HASH-BASED DATABASE IMPLEMENTATIONS

Goal / Approach

For the CPS377 Final Project we developed a C implementation of a Linear Hash Table and a B+-Tree to compare their benefits and downsides. We expected to demonstrate that B+-trees are better at range queries and hash tables are better optimized for key equality searches.

Challenges

Page simulation design

Directly collecting data on page hits was infeasible for a number of reasons, primarily time: such data would require manipulation of the Linux kernel to collect, which would have added unrealistic amounts of work to the project. Moreover, we wanted to be able to modify the page size so that we could test in small scale and benchmark across a variety of page sizes. We thus needed a way to simulate page fetching within the program itself. The solution to this is the `pageManager` struct, which is a simple set of counters for reads and writes. Pages were abstracted to a `pageptr` struct, which contains a union of the various type of pages in play (hash buckets, tree nodes, tree leaf nodes, record pages) and a signifier to denote which type that particular pointer refers to. The data structures were then implemented based on this abstracted `pageptr`, and the algorithms call the `pageManager` every time they perform an action that would require fetching a new page (e.g. going to a bucket or traversing to a tree node).

The “page” size can then be set at compile time with a pre-processing directive, with the sizes of the individual types of pages calculated based on the sizes of their component structs; so a tree node, composed of a 64-bit int and an array of key/pointer unions contains as many of those unions as fit in the page size, rounded down to an odd number. This arithmetic was the source of a number of very difficult errors to chase down, as it relies on certain assumptions about the sizes of various structs that may be invalidated based on how the C compiler packs members into a struct; this caused a number of memory errors, and so led to the implementation of a `checkSizes` function that compares the expected values against the actual at the initialization of the page manager and terminates the program if the values disagree.

Hash

One of the first challenges I encountered when writing the linear hashing data structure was how to generate these family of hashing functions. I did this by creating a function which took two parameters, a level, and the key to be hashed. Then I set all the bits to zero in the 32-bit key except for the last

level_i bits and returned the key. This was very elegant because it allowed me to return an int that could be used to directly index into the array of buckets representing the hash table.

Another challenge I encountered was when I thought I was done with the writing the data structure. To preface the description of the problem, the issue was I was never resetting the next pointer to point to the first bucket when the hash table reached the end of the level. What I noticed was the hash table was splitting *very* often and the buckets were not uniformly filled (the lower buckets were filled with lots of overflow buckets (which makes sense because these early buckets were never being split, since next never pointed to the same bucket twice). This also lead to relatively fast inserts and searches with a relatively small number of records (e.g. one-hundred thousand), but exponentially slower inserts and searches with just 10 times the number of records (e.g. one million, in fact this never finished because the process was killed because it ran out of memory). Resetting next to point to bucket 0 every time I reached the end of level, resulted in the buckets being much more uniformly filled, because there were far fewer splits. It also lead to a linear relationship between number of records inserted and the time it took to insert and search for them.

Tree

The chief difficulty in implementing the tree structure was the number of different structs that had to be managed. Not only are the individual nodes two different types depending on whether they are leaf or non-leaf nodes, but within non-leaf nodes there are interlaced keys and pointers. The former problem is actually handled very elegantly by the `pageptr` struct: the tree only needs to store `pageptr` structs instead of raw pointers, the types of which can then be checked using the signifier in the `pageptr`.

The interlaced array of keys and pointers was done with a further union of an int and a `pageptr`: the array was simply an array of that union, with every even member being a pointer and every odd a key. This could alternatively have been done using two arrays, one of keys and one of pointers; such an approach would have been perhaps more memory-efficient, but require more sophisticated size calculations.

Results

Our prediction from our project proposal (like the theory suggests) that hash tables are better optimized for key equality searches (Figure 1) and that B+-trees are better at range queries (Figure 2), was supported empirically by our benchmarked results.

Search

Figure 1 (also see Tables 1 & 2) demonstrates the performance of a key equality searches of the entire data structure for both the the linear hash and tree implementation. The hash table performs better than the tree for key equality searches. This, of course, makes sense because in key equality searches the tree cannot benefit from the links between the leaf nodes like it does in range search. The hash table performs roughly 4 times better than the tree on average. Key equality search is the worst case for tree-based databases.

Range Search

Figure 2 demonstrates the performance of a range query of the entire data structure for both the the linear hash and tree implementation. The tree performs **overwhelming** better than the hash table for range searches, the performance gap is so wide that it helps to look at the same plot but with a logarithmically scaled x-axis (Figure 3). Range search is the worst case for hash-based databases.

Operation	Reads	Writes
Insert	3223374	2857734
Full Key Equality Search	1095543	0
Second Q Range Search	17703231	0
Full Range Search	102532199	0

Table 1: Hash Table with 1,000,000 Random Skewed Records

Operation	Reads	Writes
Insert	4361765	1250899
Full Key Equality Search	6000000	0
Second Q Range Search	12368	0
Full Range Search	49310	0

Table 2: Tree with 1,000,000 Random Skewed Records




Figure 1: Performance of Search for Linear Hash Table and Tree




Figure 2: Performance of Range Search for Linear Hash Table and Tree

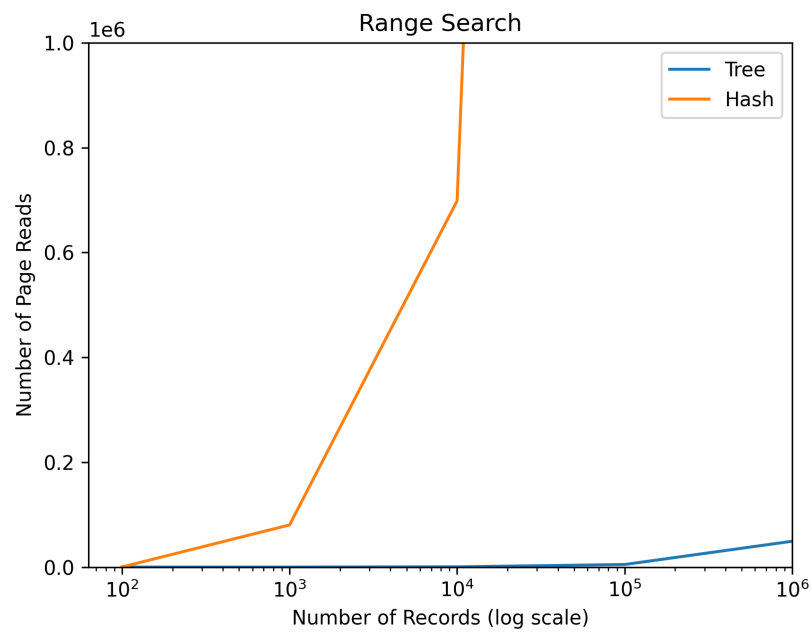


Figure 3: Performance of Range Search for Linear Hash Table and Tree with a Logarithmically Scaled X-axis